

TinyML and Efficient AI Computing MIT Course

Hariprashad Ravikumar

This document contains the lecture notes for the TinyML and Efficient AI Computing MIT course, alongside older personal machine learning notes and mathematical derivations.

Contents

1	TinyML and Efficient AI Computing	2
1.1	Lecture 2 - Basics of Neural Networks	2
2	Hari's personal note on Neural Networks	3

1 TinyML and Efficient AI Computing

1.1 Lecture 2 - Basics of Neural Networks

Deep Neural Network

3-Layer Neural Network With 2 Hidden Layers

Inputs

Layer 0

Hidden Layer 1

Hidden Layer 2

Outputs

The dimensionality of these *hidden* layers determines the **width** of the model.

Synapses / Weights / Parameters

Neurons / Features / Activations

Synapse

Input Axon

Dendrite

Cell Body

Output Axon

Activation Function

$y_j = f\left(\sum_i w_i x_i + b\right)$

MIT 6.5940: TinyML and Efficient Deep Learning Computing <https://efficientml.ai> 11

Fully-Connected Layer (Linear Layer)

The output neuron is connected to all input neurons.

- **Shape of Tensors:**
 - Input Features $\mathbf{X} : (1, c_i)$
 - Output Features $\mathbf{Y} : (1, c_o)$
 - Weights $\mathbf{W} : (c_o, c_i)$
 - Bias $\mathbf{b} : (c_o,)$

$$y_i = \sum_j w_{ij} x_j + b_i$$

Notations	
c_i	Input Channels
c_o	Output Channels

$$\begin{matrix} c_i \\ \boxed{} \end{matrix} \times \begin{matrix} c_o \\ \boxed{} \\ c_i \end{matrix} = \begin{matrix} c_o \\ \boxed{} \end{matrix}$$

$\mathbf{X} \quad \mathbf{W}^T \quad \mathbf{Y}$

MIT 6.5940: TinyML and Efficient Deep Learning Computing <https://efficientml.ai> 13

2 Hari's personal note on Neural Networks

2.1 A Simple Neural Network




Figure 1: Simple Two Layer Neural Network (NN)

- m = number of training examples.
- $X \in \mathbb{R}^{n_x \times m}$: input data, each column $x^{(i)}$ is one example.
- $Y \in \mathbb{R}^{n_y \times m}$: output data
- For example, in MNIST, each image is $28 \times 28 = 784$ pixels; we want have output of 10 row vector $(0, 1, 2, \dots, 9)$. In this case, $n_x = 784$, $n_y = 10$. For simplicity let's choose the number of nodes in the hidden layer as $n_h = 10$

$$A^{[0]} = X \in \mathbb{R}^{784 \times m} \quad (1)$$

$$W^{[1]} \in \mathbb{R}^{10 \times 784} \quad (2)$$

$$b^{[1]} \in \mathbb{R}^{10 \times m} \quad (3)$$

- **Layer 1 (hidden):**

$$W^{[1]} \in \mathbb{R}^{n_h \times n_x}, \quad b^{[1]} \in \mathbb{R}^{n_h \times m},$$

$$Z^{[1]} = W^{[1]} X + b^{[1]}, \quad A^{[1]} = g(Z^{[1]}),$$

where $g(\cdot)$ is an element-wise activation (e.g., ReLU or tanh).

- **Layer 2 (output):**

$$W^{[2]} \in \mathbb{R}^{n_y \times n_h}, \quad b^{[2]} \in \mathbb{R}^{n_y \times m},$$

$$Z^{[2]} = W^{[2]} A^{[1]} + b^{[2]}, \quad A^{[2]} = \text{softmax}(Z^{[2]}),$$

so that each column $A_{:,i}^{[2]}$ is a probability vector over n_y classes for example i .

We train with the average cross-entropy loss:

$$J = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^{n_y} Y_{k,i} \log(A_{k,i}^{[2]}). \quad (4)$$

Our goal is to compute the gradients

$$\frac{\partial J}{\partial W^{[2]}}, \quad \frac{\partial J}{\partial b^{[2]}}, \quad \frac{\partial J}{\partial W^{[1]}}, \quad \frac{\partial J}{\partial b^{[1]}},$$

so that we can update parameters by (for example) gradient descent. Below is the full chain-rule derivation.

2.2 Forward Propagation

2.2.1 Input layer $\rightarrow$ Hidden layer

$$Z^{[1]} = W^{[1]} X + b^{[1]} \in \mathbb{R}^{n_h \times m}, \quad A^{[1]} = g(Z^{[1]}) \in \mathbb{R}^{n_h \times m}.$$

Here g is applied element-wise (e.g., ReLU: $g(z) = \max(0, z)$, or $\tanh(z)$, etc.).

2.2.2 Hidden layer $\rightarrow$ Output layer

$$Z^{[2]} = W^{[2]} A^{[1]} + b^{[2]} \in \mathbb{R}^{n_y \times m}, \quad A^{[2]} = \text{softmax}(Z^{[2]}) \in \mathbb{R}^{n_y \times m}.$$

Concretely, for each example i and each class k ,

$$A_{k,i}^{[2]} = \frac{\exp(Z_{k,i}^{[2]})}{\sum_{j=1}^{n_y} \exp(Z_{j,i}^{[2]})}, \quad k = 1, \dots, n_y.$$

2.2.3 Compute loss

$$J = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^{n_y} Y_{k,i} \log(A_{k,i}^{[2]}).$$

2.3 Backpropagation: Output-Layer Gradients

Because we use softmax plus cross-entropy, there is a well-known simplification:

$$\frac{\partial J}{\partial Z^{[2]}} = A^{[2]} - Y \in \mathbb{R}^{n_y \times m}. \quad (5)$$

More explicitly, for each class k and example i ,

$$[dZ^{[2]}]_{k,i} = A_{k,i}^{[2]} - Y_{k,i}.$$

We denote this matrix by

$$dZ^{[2]} \equiv A^{[2]} - Y.$$

2.3.1 Gradient w.r.t. $W^{[2]}$

From

$$Z^{[2]} = W^{[2]} A^{[1]} + b^{[2]},$$

treating $A^{[1]}$ as constant,

$$\frac{\partial J}{\partial W^{[2]}} = \frac{\partial J}{\partial Z^{[2]}} \frac{\partial Z^{[2]}}{\partial W^{[2]}}.$$

In matrix form, since $Z^{[2]} = W^{[2]} A^{[1]} + b^{[2]}$,

$$\frac{\partial J}{\partial W^{[2]}} = \frac{1}{m} (dZ^{[2]}) (A^{[1]})^T \in \mathbb{R}^{n_y \times n_h}. \quad (6)$$

(The factor $1/m$ comes from averaging over m examples.)

2.3.2 Gradient w.r.t. $b^{[2]}$

Because $b^{[2]}$ is broadcast-added to each column of $W^{[2]} A^{[1]}$,

$$\frac{\partial J}{\partial b^{[2]}} = \frac{1}{m} \sum_{i=1}^m [dZ^{[2]}]_{:,i} = \frac{1}{m} dZ^{[2]} \mathbf{1}_m,$$

where $\mathbf{1}_m \in \mathbb{R}^m$ is a column-vector of all ones. In practice:

$$db^{[2]} = \frac{1}{m} \sum_{i=1}^m dZ_{:,i}^{[2]} \in \mathbb{R}^{n_y \times 1}. \quad (7)$$

2.3.3 Gradient w.r.t. $A^{[1]}$ (propagate to previous layer)

Since $Z^{[2]} = W^{[2]} A^{[1]} + b^{[2]}$,

$$dA^{[1]} = \frac{\partial J}{\partial A^{[1]}} = (W^{[2]})^T dZ^{[2]} \in \mathbb{R}^{n_h \times m}. \quad (8)$$

We often call this intermediate term $dA^{[1]} = (W^{[2]})^T dZ^{[2]}$.

2.4 Backpropagation: Hidden-Layer Gradients

At layer1, we have

$$Z^{[1]} = W^{[1]} X + b^{[1]}, \quad A^{[1]} = g(Z^{[1]}),$$

where g is the activation (e.g., ReLU or tanh). We already computed

$$dA^{[1]} = (W^{[2]})^T dZ^{[2]} \in \mathbb{R}^{n_h \times m}.$$

2.4.1 Compute $dZ^{[1]}$

By the chain rule,

$$dZ^{[1]} = \frac{\partial J}{\partial Z^{[1]}} = dA^{[1]} \circ g'(Z^{[1]}), \quad (9)$$

where “ $\circ$ ” denotes element-wise (Hadamard) multiplication, and

$$g'(Z^{[1]}) = \text{matrix of derivatives of } g \text{ evaluated at each entry of } Z^{[1]}.$$

- If $g = \text{ReLU}$, then

$$g'(z) = \begin{cases} 1, & z > 0, \\ 0, & z \leq 0, \end{cases} \implies [g'(Z^{[1]})]_{j,i} = \begin{cases} 1, & Z_{j,i}^{[1]} > 0, \\ 0, & Z_{j,i}^{[1]} \leq 0. \end{cases}$$

- If $g = \tanh$, then

$$g'(z) = 1 - \tanh^2(z), \implies [g'(Z^{[1]})]_{j,i} = 1 - (A_{j,i}^{[1]})^2.$$

- If $g = \sigma$ (sigmoid), then $g'(z) = \sigma(z)(1 - \sigma(z))$.

Thus

$$dZ^{[1]} = dA^{[1]} \circ g'(Z^{[1]}) \in \mathbb{R}^{n_h \times m}.$$

2.4.2 Gradient w.r.t. $W^{[1]}$

From

$$Z^{[1]} = W^{[1]} X + b^{[1]},$$

treating X as constant,

$$\frac{\partial J}{\partial W^{[1]}} = \frac{\partial J}{\partial Z^{[1]}} \frac{\partial Z^{[1]}}{\partial W^{[1]}} = \frac{1}{m} (dZ^{[1]}) X^T \in \mathbb{R}^{n_h \times n_x}.$$

Hence

$$dW^{[1]} = \frac{1}{m} dZ^{[1]} X^T. \quad (10)$$

2.4.3 Gradient w.r.t. $b^{[1]}$

Because $b^{[1]}$ is added to each column of $W^{[1]}X$,

$$\frac{\partial J}{\partial b^{[1]}} = \frac{1}{m} \sum_{i=1}^m [dZ^{[1]}]_{:,i} \in \mathbb{R}^{n_h \times 1}.$$

Thus

$$db^{[1]} = \frac{1}{m} \sum_{i=1}^m dZ_{:,i}^{[1]}. \quad (11)$$

2.5 Full Overview

Putting it all together for a two-layer network (with a single hidden layer of size n_h):

2.5.1 Forward Propagation

$$\begin{cases} Z^{[1]} = W^{[1]} X + b^{[1]}, & A^{[1]} = g(Z^{[1]}), \\ Z^{[2]} = W^{[2]} A^{[1]} + b^{[2]}, & A^{[2]} = \text{softmax}(Z^{[2]}), \end{cases}$$

$$J = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^{n_y} Y_{k,i} \log(A_{k,i}^{[2]}).$$

2.5.2 Backward Propagation

Output layer

$$dZ^{[2]} = A^{[2]} - Y \in \mathbb{R}^{n_y \times m}, \quad (12)$$

$$dW^{[2]} = \frac{1}{m} dZ^{[2]} (A^{[1]})^T \in \mathbb{R}^{n_y \times n_h}, \quad (13)$$

$$db^{[2]} = \frac{1}{m} \sum_{i=1}^m dZ^{[2]}_{:,i} \in \mathbb{R}^{n_y \times 1}, \quad (14)$$

$$dA^{[1]} = (W^{[2]})^T dZ^{[2]} \in \mathbb{R}^{n_h \times m}. \quad (15)$$

Hidden layer

$$dZ^{[1]} = dA^{[1]} \circ g'(Z^{[1]}) \in \mathbb{R}^{n_h \times m}, \quad (16)$$

$$dW^{[1]} = \frac{1}{m} dZ^{[1]} X^T \in \mathbb{R}^{n_h \times n_x}, \quad (17)$$

$$db^{[1]} = \frac{1}{m} \sum_{i=1}^m dZ^{[1]}_{:,i} \in \mathbb{R}^{n_h \times 1}. \quad (18)$$

2.5.3 Parameter Updates (Gradient Descent)

After computing these gradients, update each parameter by

$$W^{[\ell]} \leftarrow W^{[\ell]} - \alpha dW^{[\ell]}, \quad b^{[\ell]} \leftarrow b^{[\ell]} - \alpha db^{[\ell]}, \quad \ell = 1, 2,$$

where α is the learning rate.

2.6 Generalization to More Layers

If there are L layers, the pattern is:

$$Z^{[\ell]} = W^{[\ell]} A^{[\ell-1]} + b^{[\ell]}, \quad A^{[\ell]} = g^{[\ell]}(Z^{[\ell]}), \quad \ell = 1, \dots, L,$$

with $A^{[0]} = X$. The final layer L uses softmax (for multiclass) or sigmoid (for binary).

During backpropagation:

Compute $dZ^{[L]}$ from the loss derivative at the output (e.g., $A^{[L]} - Y$).

For $\ell = L, L-1, \dots, 2$:

$$dW^{[\ell]} = \frac{1}{m} dZ^{[\ell]} (A^{[\ell-1]})^T,$$

$$db^{[\ell]} = \frac{1}{m} \sum_{i=1}^m dZ^{[\ell]}_{:,i},$$

$$dA^{[\ell-1]} = (W^{[\ell]})^T dZ^{[\ell]},$$

$$dZ^{[\ell-1]} = dA^{[\ell-1]} \circ g'^{[\ell-1]}(Z^{[\ell-1]}).$$

Then update each $W^{[\ell]}$, $b^{[\ell]}$ accordingly.

References

Michael A. Nielsen, *Neural Networks and Deep Learning*, Determination Press, 2015.